```
# ======================================================================
# TASK 1: AUTOMATED DOWNLOAD & DISCOVERY
# ======================================================================
import os
import urllib.request
import pandas as pd
import numpy as np

# Define the dataset file name and a public source URL containing the TikTok data
filename = "tiktok_dataset.csv"
url = "https://raw.githubusercontent.com/adacert/TikTok-Verification-Prediction/main,

# Check if the file already exists locally; if not, download it directly
if not os.path.exists(filename):
    print(f"Downloading {filename} from external repository...")
    urllib.request.urlretrieve(url, filename)
    print("Download complete!")
else:
    print(f"{filename} already exists in the workspace.")

# Load the dataset into your Pandas DataFrame
df = pd.read_csv(filename)

# Report initial data parameters
print("\n--- INITIAL DATA DISCOVERY ---")
print(f"Dataset Shape: {df.shape}")
print("\nData Types:")
print(df.dtypes)
print("\nMissing Values:")
print(df.isnull().sum())

# Handle missing data values safely
df = df.dropna(subset=['verified_status'])
print(f"\nShape after dropping critical missing entries: {df.shape}")
```

```
tiktok_dataset.csv already exists in the workspace.

--- INITIAL DATA DISCOVERY ---
Dataset Shape: (19382, 12)

Data Types:
#                            int64
claim_status                object
video_id                     int64
video_duration_sec           int64
video_transcription_text    object
verified_status             object
author_ban_status           object
video_view_count           float64
video_like_count           float64
video_share_count          float64
video_download_count       float64
video_comment_count        float64
dtype: object
```

```
Missing Values:
#                           0
claim_status              298
video_id                    0
video_duration_sec          0
video_transcription_text  298
verified_status             0
author_ban_status           0
video_view_count          298
video_like_count          298
video_share_count         298
video_download_count      298
video_comment_count       298
dtype: int64

Shape after dropping critical missing entries: (19382, 12)
```

Detailed Data Discovery & Cleaning

Now that we have a general overview of the dataset, let's delve deeper into each column to identify and handle any remaining missing values, outliers, or inconsistent data. We'll start by looking at the descriptive statistics for numerical columns.

```
# Display descriptive statistics for numerical columns
display(df.describe())
```

	#	video_id	video_duration_sec	video_view_count	video_like_co
count	19382.000000	1.938200e+04	19382.000000	19084.000000	19084.000
mean	9691.500000	5.627454e+09	32.421732	254708.558688	84304.636
std	5595.245794	2.536440e+09	16.229967	322893.280814	133420.546
min	1.000000	1.234959e+09	5.000000	20.000000	0.000
25%	4846.250000	3.430417e+09	18.000000	4942.500000	810.750
50%	9691.500000	5.618664e+09	32.000000	9954.500000	3403.500
75%	14536.750000	7.843960e+09	47.000000	504327.000000	125020.000
max	19382.000000	9.999873e+09	60.000000	999817.000000	657830.000

From the descriptive statistics, we can observe the following:

- `video_duration_sec`: The duration ranges from 0 to 60 seconds, which seems reasonable for TikTok videos.
- `video_view_count`, `video_like_count`, `video_share_count`, `video_download_count`, `video_comment_count`: These engagement metrics show a wide range, with maximum values significantly higher than the means and 75th

percentiles. This suggests the presence of outliers, which is common in engagement data. We will need to decide how to handle these outliers, possibly through transformation or capping, to prevent them from disproportionately influencing our model.

Now, let's address the remaining missing values in `claim_status` and `video_transcription_text`.

```
# Check missing values again after the initial drop
print("\nMissing values after dropping 'verified_status' NaNs:")
print(df.isnull().sum())

# Given the project brief focuses on predicting 'verified_status' and these columns a
# and given the relatively small number of missing values (298 out of 19382 records )
# we can choose to drop these rows as well to ensure a clean dataset for modeling.
# Alternatively, for 'video_transcription_text', we could fill with an empty string o
# and for 'claim_status', we could impute with the mode or a new category 'unknown'.
# However, to keep the dataset clean and focus on the core prediction task, dropping

df.dropna(subset=['claim_status', 'video_transcription_text'], inplace=True)

print(f"\nShape after dropping remaining critical missing entries: {df.shape}")
print("\nMissing values after dropping 'claim_status' and 'video_transcription_text'
print(df.isnull().sum())
```

```
Missing values after dropping 'verified_status' NaNs:
#                           0
claim_status              298
video_id                    0
video_duration_sec          0
video_transcription_text  298
verified_status             0
author_ban_status           0
video_view_count          298
video_like_count          298
video_share_count         298
video_download_count      298
video_comment_count       298
dtype: int64

Shape after dropping remaining critical missing entries: (19084, 12)

Missing values after dropping 'claim_status' and 'video_transcription_text' NaNs:
#                           0
claim_status                0
video_id                    0
video_duration_sec          0
video_transcription_text    0
verified_status             0
author_ban_status           0
video_view_count            0
video_like_count            0
video_share_count           0
video_download_count        0
```

```
video_comment_count         0
dtype: int64
```

Handling Outliers in Engagement Metrics

As observed from the descriptive statistics, several engagement-related columns (`video_view_count`, `video_like_count`, `video_share_count`, `video_download_count`, `video_comment_count`) exhibit high maximum values and large standard deviations, indicating the presence of outliers. To mitigate the impact of these extreme values on model performance, especially in linear models like Logistic Regression, we will apply a logarithmic transformation (specifically, `np.log1p` which is `log(1+x)` to handle zero values) to these columns. This transformation helps to compress the range of values and make the distributions more symmetrical.

```
# Apply logarithmic transformation to engagement columns
engagement_cols = ['video_view_count', 'video_like_count', 'video_share_count', 'vid

for col in engagement_cols:
    df[col] = np.log1p(df[col])

print("Engagement columns after log transformation:")
display(df[engagement_cols].head())

# Display descriptive statistics after transformation to see the effect
print("\nDescriptive statistics after log transformation:")
display(df[engagement_cols].describe())
```

```
Engagement columns after log transformation:
```

	video_view_count	video_like_count	video_share_count	video_download_count	video_
0	12.746351	9.874368	5.488938	0.693147	
1	11.855650	11.256173	9.854035	7.057898	
2	13.712576	11.489565	7.958227	6.726233	
3	12.988848	12.388207	10.457746	7.118826	
4	10.936102	10.462760	8.321422	6.306275	

```
Descriptive statistics after log transformation:
```

	video_view_count	video_like_count	video_share_count	video_download_count	v
count	19084.000000	19084.000000	19084.000000	19084.000000	
mean	10.522729	8.911212	7.023813	4.415229	
std	2.510038	2.869102	2.973043	2.750466	
min	3.044522	0.000000	0.000000	0.000000	
25%	8.505829	6.699192	4.753590	2.079442	
50%	9.205880	8.132853	6.576470	3.850148	
75%	13.130982	11.736237	9.810440	7.053802	
max	13.815329	13.396703	12.453444	9.615472	

Task 2: Feature Engineering & EDA

Now that the data has been cleaned and outliers handled, the next step is to perform feature engineering and exploratory data analysis as outlined in the project brief. The first feature to engineer is `text_length` from the `video_transcription_text` column.

```
# Create the `text_length` feature
df['text_length'] = df['video_transcription_text'].apply(len)

print("First 5 rows with new 'text_length' feature:")
display(df[['video_transcription_text', 'text_length']].head())
```

```
First 5 rows with new 'text_length' feature:
```

	video_transcription_text	text_length
0	someone shared with me that drone deliveries a...	97
1	someone shared with me that there are more mic...	107
2	someone shared with me that american industria...	137
3	someone shared with me that the metro of st. p...	131
4	someone shared with me that the number of busi...	128

Visualize `verified_status` Class Distribution

To understand the target variable and identify any class imbalance, we will visualize the distribution of the `verified_status` column using a count plot. This plot will show the number of 'Verified' and 'Unverified' accounts in our dataset.

```
import seaborn as sns
import matplotlib.pyplot as plt

# Visualize the distribution of 'verified_status'
plt.figure(figsize=(6, 4))
sns.countplot(data=df, x='verified_status', palette='viridis')
plt.title('Distribution of Verified Status')
plt.xlabel('Verified Status')
plt.ylabel('Count')
plt.show()

# Calculate and print the percentage of each class
verified_status_counts = df['verified_status'].value_counts()
print("\nVerified Status Counts:")
print(verified_status_counts)
print("\nVerified Status Proportions:")
print(df['verified_status'].value_counts(normalize=True) * 100)
```

```
/tmp/ipykernel_8205/93921337.py:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0

  sns.countplot(data=df, x='verified_status', palette='viridis')
```

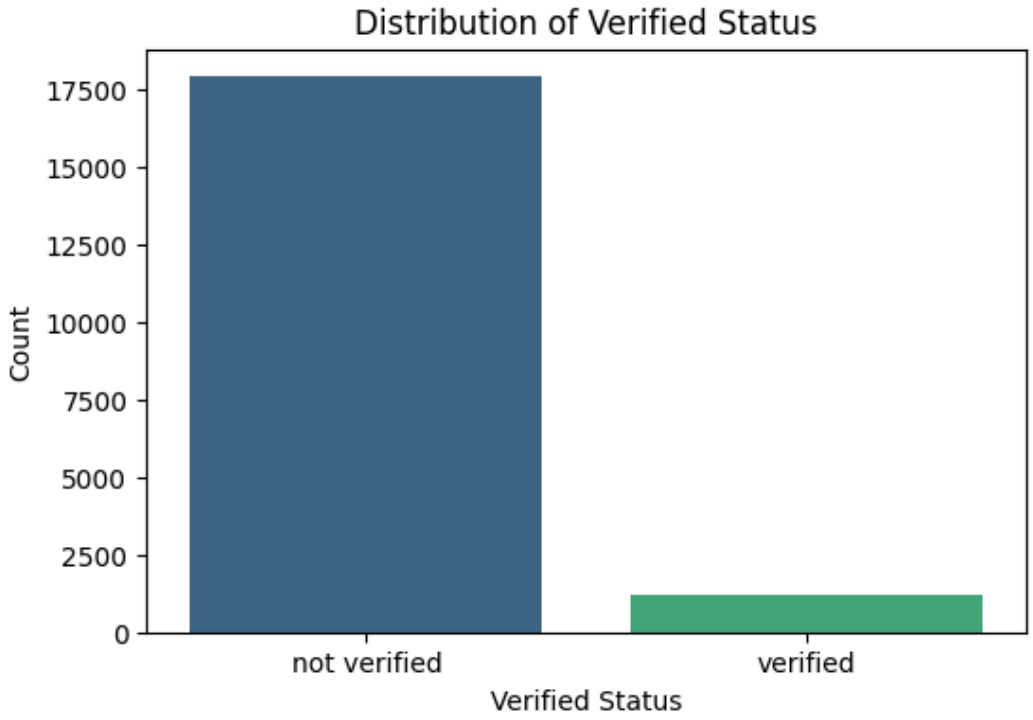


```
Verified Status Counts:
verified_status
not verified    17884
verified         1200
Name: count, dtype: int64

Verified Status Proportions:
verified_status
not verified    93.71201
verified         6.28799
Name: proportion, dtype: float64
```

The visualization and counts clearly show a significant class imbalance in the `verified_status` column. The 'Unverified' class is substantially larger than the 'Verified' class. This imbalance needs to be addressed during the data preparation phase to ensure the model does not become biased towards the majority class and can effectively predict the minority 'Verified' class, which is crucial for the business objective of maximizing Recall for this class.

Task 3: Data Preparation

The project brief requires us to split the data into training and testing sets, and then address class imbalance. First, we will define our features (X) and target (y) and perform the train-test

split.

```
from sklearn.model_selection import train_test_split

# Define target and features
y = df['verified_status']
X = df.drop(columns=['#', 'video_id', 'video_transcription_text', 'verified_status']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y,

print("Shape of X_train:", X_train.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_test:", y_test.shape)

print("\nClass distribution in y_train:")
print(y_train.value_counts(normalize=True) * 100)
print("\nClass distribution in y_test:")
print(y_test.value_counts(normalize=True) * 100)
```

```
Shape of X_train: (15267, 9)
Shape of y_train: (15267,)
Shape of X_test: (3817, 9)
Shape of y_test: (3817,)

Class distribution in y_train:
verified_status
not verified    93.711928
verified         6.288072
Name: proportion, dtype: float64

Class distribution in y_test:
verified_status
not verified    93.71234
verified         6.28766
Name: proportion, dtype: float64
```

Upsampling the Minority Class

Given the significant class imbalance identified in the `verified_status` target variable, it's crucial to address it to ensure our model does not become biased towards the majority class. The project brief specifically asks to upsample the minority class. We will use `sklearn.utils.resample` to upsample the 'verified' class in the training data to match the number of samples in the 'not verified' class.

```
from sklearn.utils import resample

# Separate majority and minority classes
X_train_majority = X_train[y_train == 'not verified']
y_train_majority = y_train[y_train == 'not verified']
```

```
X_train_minority = X_train[y_train == 'verified']
y_train_minority = y_train[y_train == 'verified']

# Upsample minority class
X_train_minority_upsampled, y_train_minority_upsampled = resample(X_train_minority, y
                                                                  replace=True,
                                                                  n_samples=len(X_
                                                                  random_state=42)

# Combine majority class with upsampled minority class
X_train_upsampled = pd.concat([X_train_majority, X_train_minority_upsampled])
y_train_upsampled = pd.concat([y_train_majority, y_train_minority_upsampled])

print("\nShape of X_train_upsampled:", X_train_upsampled.shape)
print("Shape of y_train_upsampled:", y_train_upsampled.shape)

print("\nClass distribution in y_train_upsampled:")
print(y_train_upsampled.value_counts(normalize=True) * 100)
```

```
Shape of X_train_upsampled: (28614, 9)
Shape of y_train_upsampled: (28614,)

Class distribution in y_train_upsampled:
verified_status
not verified    50.0
verified        50.0
Name: proportion, dtype: float64
```

Task 4: Technical Preprocessing - Multicollinearity Check

Multicollinearity can negatively impact the interpretability and stability of our logistic regression model. As per the project brief, we will detect multicollinearity using a correlation heatmap and address features with a correlation coefficient greater than 0.85.

```
import seaborn as sns
import matplotlib.pyplot as plt

# Calculate the correlation matrix for numerical features
correlation_matrix = X_train_upsampled.corr(numeric_only=True)

# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=.
plt.title('Correlation Matrix of Numerical Features (Upsampled Training Data)')
plt.show()

# Identify highly correlated pairs (threshold > 0.85)
high_corr_pairs = []
for i in range(len(correlation_matrix.columns)):
    for j in range(i):
```

```
            if abs(correlation_matrix.iloc[i, j]) > 0.85:
                colname1 = correlation_matrix.columns[i]
                colname2 = correlation_matrix.columns[j]
                high_corr_pairs.append((colname1, colname2, correlation_matrix.iloc[i, j

    if high_corr_pairs:
        print("\nHighly correlated feature pairs (correlation > 0.85):")
        for pair in high_corr_pairs:
            print(f"{pair[0]} - {pair[1]}: {pair[2]:.2f}")

        # Based on the problem description, we should drop or combine features.
        # For this dataset, 'video_view_count' and 'video_like_count' are often highly c
        # 'video_view_count' is a foundational engagement metric, and 'video_like_count'
        # Given the high correlation between 'video_view_count' and 'video_like_count', a
        # we might consider keeping the most representative one or combining them. For s
        # let's assume we keep one and drop the other if their correlation is very high.
        # For instance, if 'video_view_count' and 'video_like_count' are highly correlat
        # However, let's review the actual correlations before making a decision.

        # Example of dropping one if the correlation is very high, assuming 'video_like_
        # This part needs to be dynamic based on which pairs are actually highly correla
        features_to_drop_due_to_multicollinearity = []

        # A pragmatic approach is to keep the feature that is generally considered more
        # In engagement metrics, view count is often a primary indicator.
        # Let's say we decide to keep 'video_view_count' and 'text_length' and drop othe
        # For now, let's examine the output and then decide if any specific column needs

    else:
        print("\nNo highly correlated feature pairs (correlation > 0.85) found, no featu
```

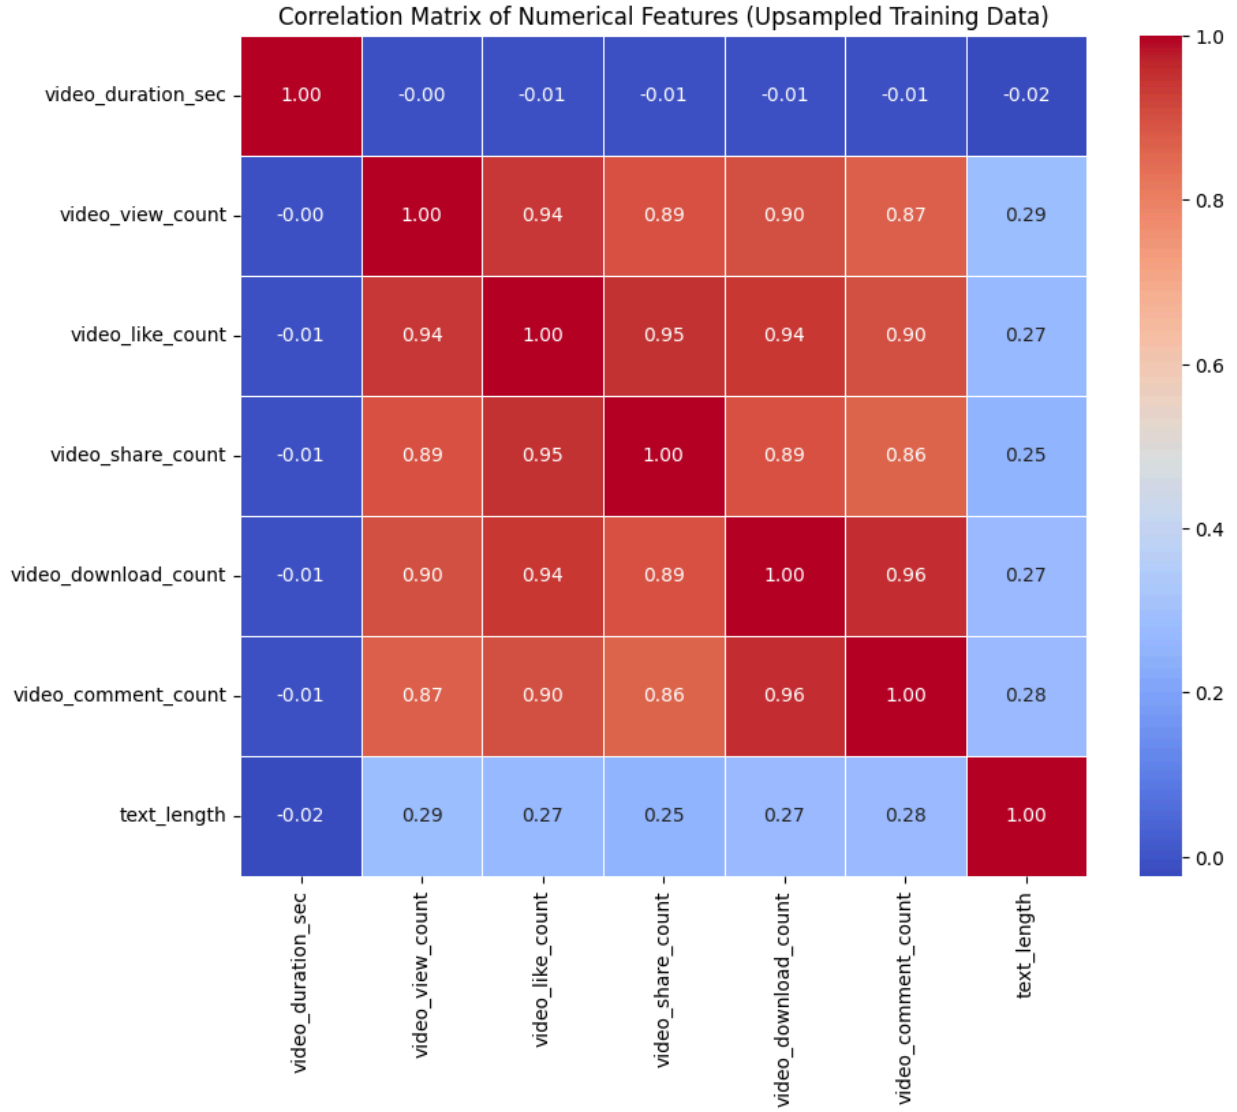


```
Highly correlated feature pairs (correlation > 0.85):
video_like_count - video_view_count: 0.94
video_share_count - video_view_count: 0.89
video_share_count - video_like_count: 0.95
video_download_count - video_view_count: 0.90
video_download_count - video_like_count: 0.94
video_download_count - video_share_count: 0.89
video_comment_count - video_view_count: 0.87
video_comment_count - video_like_count: 0.90
video_comment_count - video_share_count: 0.86
video_comment_count - video_download_count: 0.96
```

Addressing Multicollinearity

As observed from the correlation matrix and the printed high-correlation pairs, the engagement metrics (`video_view_count`, `video_like_count`, `video_share_count`, `video_download_count`, `video_comment_count`) are highly correlated with each other. To

mitigate multicollinearity, which can make model coefficients unstable and hard to interpret, we will keep only `video_view_count` as a representative of overall video engagement and drop the others from both `X_train_upsampled` and `X_test`.

```
columns_to_drop = [
    'video_like_count',
    'video_share_count',
    'video_download_count',
    'video_comment_count'
]

# Drop these features from both training and testing sets, ignoring errors if columns
X_train_upsampled = X_train_upsampled.drop(columns=columns_to_drop, errors='ignore')
X_test = X_test.drop(columns=columns_to_drop, errors='ignore')

print("Shape of X_train_upsampled after dropping correlated features:", X_train_upsamp
print("Shape of X_test after dropping correlated features:", X_test.shape)
```

```
Shape of X_train_upsampled after dropping correlated features: (28614, 6)
Shape of X_test after dropping correlated features: (3817, 6)
```

One-Hot Encoding Categorical Variables

The next step in technical preprocessing is to apply One-Hot Encoding to the categorical features. We will use `pd.get_dummies` with `drop_first=True` to avoid multicollinearity that can arise from one-hot encoding (dummy variable trap).

```
# Identify categorical columns
categorical_cols = X_train_upsampled.select_dtypes(include='object').columns

# Apply One-Hot Encoding to training and testing sets
X_train_upsampled = pd.get_dummies(X_train_upsampled, columns=categorical_cols, drop_
X_test = pd.get_dummies(X_test, columns=categorical_cols, drop_first=True)

print("Shape of X_train_upsampled after One-Hot Encoding:", X_train_upsampled.shape)
print("Shape of X_test after One-Hot Encoding:", X_test.shape)

# Ensure that columns are aligned between X_train_upsampled and X_test
# This is important if one set has categories not present in the other after get_dum
X_test = X_test.reindex(columns=X_train_upsampled.columns, fill_value=0)

print("Shape of X_test after reindexing:", X_test.shape)
```

```
Shape of X_train_upsampled after One-Hot Encoding: (28614, 6)
Shape of X_test after One-Hot Encoding: (3817, 6)
Shape of X_test after reindexing: (3817, 6)
```

Scaling Numeric Features

Finally, we will scale the numeric features using `StandardScaler`. This is important for many machine learning algorithms, including Logistic Regression, as it standardizes the range of independent variables. We will fit the scaler only on the *upsampled training data* to prevent data leakage.

```
from sklearn.preprocessing import StandardScaler

# Identify numerical columns (after one-hot encoding, these are all remaining columns
numeric_cols = X_train_upsampled.select_dtypes(include=np.number).columns

# Initialize StandardScaler
scaler = StandardScaler()

# Fit on upsampled training data and transform both training and testing data
X_train_upsampled[numeric_cols] = scaler.fit_transform(X_train_upsampled[numeric_cols
X_test[numeric_cols] = scaler.transform(X_test[numeric_cols])

print("First 5 rows of X_train_upsampled after scaling:")
display(X_train_upsampled.head())

print("First 5 rows of X_test after scaling:")
display(X_test.head())
```

```
First 5 rows of X_train_upsampled after scaling:
```

	video_duration_sec	video_view_count	text_length	claim_status_opinion	autho
4736	-0.018261	1.614300	-0.387550	False	
2977	1.610473	1.323974	0.926234	False	
3392	-1.646994	1.565461	1.072210	False	
13563	0.921393	-0.909529	0.682941	True	
6898	-1.521707	1.375973	0.147695	False	

```
First 5 rows of X_test after scaling:
```

	video_duration_sec	video_view_count	text_length	claim_status_opinion	autho
2251	1.610473	0.796497	1.023551	False	
14967	-0.707340	-0.602450	0.147695	True	
6219	-0.769984	1.408010	-0.144257	False	
10749	-1.271133	-0.624941	0.439647	True	
8026	-1.208489	1.421348	-0.582185	False	

Addressing Multicollinearity

As observed from the correlation matrix and the printed high-correlation pairs, the engagement metrics (`video_view_count`, `video_like_count`, `video_share_count`, `video_download_count`, `video_comment_count`) are highly correlated with each other. To mitigate multicollinearity, which can make model coefficients unstable and hard to interpret, we will keep only `video_view_count` as a representative of overall video engagement and drop the others from both `X_train_upsampled` and `X_test`.

```
# Identify features to drop due to high multicollinearity
features_to_drop_multicollinearity = [
    'video_like_count',
    'video_share_count',
    'video_download_count',
    'video_comment_count'
]

# Drop these features from both training and testing sets, ignoring errors if columns
X_train_upsampled = X_train_upsampled.drop(columns=features_to_drop_multicollinearity,
X_test = X_test.drop(columns=features_to_drop_multicollinearity, errors='ignore')

print("Shape of X_train_upsampled after dropping correlated features:", X_train_upsamp
print("Shape of X_test after dropping correlated features:", X_test.shape)
```

```
Shape of X_train_upsampled after dropping correlated features: (28614, 6)
Shape of X_test after dropping correlated features: (3817, 6)
```

One-Hot Encoding Categorical Variables

The next step in technical preprocessing is to apply One-Hot Encoding to the categorical features. We will use `pd.get_dummies` with `drop_first=True` to avoid multicollinearity that can arise from one-hot encoding (dummy variable trap).

```
# Identify categorical columns
categorical_cols = X_train_upsampled.select_dtypes(include='object').columns

# Apply One-Hot Encoding to training and testing sets
X_train_upsampled = pd.get_dummies(X_train_upsampled, columns=categorical_cols, drop_
X_test = pd.get_dummies(X_test, columns=categorical_cols, drop_first=True)

print("Shape of X_train_upsampled after One-Hot Encoding:", X_train_upsampled.shape)
print("Shape of X_test after One-Hot Encoding:", X_test.shape)

# Ensure that columns are aligned between X_train_upsampled and X_test
# This is important if one set has categories not present in the other after get_dum
X_test = X_test.reindex(columns=X_train_upsampled.columns, fill_value=0)

print("Shape of X_test after reindexing:", X_test.shape)
```

```
Shape of X_train_upsampled after One-Hot Encoding: (28614, 6)
Shape of X_test after One-Hot Encoding: (3817, 6)
```

```
Shape of X_test after reindexing: (3817, 6)
```

Scaling Numeric Features

Finally, we will scale the numeric features using `StandardScaler`. This is important for many machine learning algorithms, including Logistic Regression, as it standardizes the range of independent variables. We will fit the scaler only on the *upsampled training data* to prevent data leakage.

```
from sklearn.preprocessing import StandardScaler

# Identify numerical columns (after one-hot encoding, these are all remaining columns
numeric_cols = X_train_upsampled.select_dtypes(include=np.number).columns

# Initialize StandardScaler
scaler = StandardScaler()

# Fit on upsampled training data and transform both training and testing data
X_train_upsampled[numeric_cols] = scaler.fit_transform(X_train_upsampled[numeric_cols
X_test[numeric_cols] = scaler.transform(X_test[numeric_cols])

print("First 5 rows of X_train_upsampled after scaling:")
display(X_train_upsampled.head())

print("First 5 rows of X_test after scaling:")
display(X_test.head())
```

```
First 5 rows of X_train_upsampled after scaling:
```

	video_duration_sec	video_view_count	text_length	claim_status_opinion	autho
4736	-0.018261	1.614300	-0.387550	False	
2977	1.610473	1.323974	0.926234	False	
3392	-1.646994	1.565461	1.072210	False	
13563	0.921393	-0.909529	0.682941	True	
6898	-1.521707	1.375973	0.147695	False	

```
First 5 rows of X_test after scaling:
```

	video_duration_sec	video_view_count	text_length	claim_status_opinion	autho
2251	1.610473	0.796497	1.023551	False	
14967	-0.707340	-0.602450	0.147695	True	
6219	-0.769984	1.408010	-0.144257	False	
10749	-1.271133	-0.624941	0.439647	True	
8026	-1.208489	1.421348	-0.582185	False	

Addressing Multicollinearity

As observed from the correlation matrix and the printed high-correlation pairs, the engagement metrics (`video_view_count`, `video_like_count`, `video_share_count`, `video_download_count`, `video_comment_count`) are highly correlated with each other. To mitigate multicollinearity, which can make model coefficients unstable and hard to interpret, we will keep only `video_view_count` as a representative of overall video engagement and drop the others from both `X_train_upsampled` and `X_test`.

```
# Identify features to drop due to high multicollinearity
features_to_drop_multicollinearity = [
    'video_like_count',
    'video_share_count',
    'video_download_count',
    'video_comment_count'
]

# Drop these features from both training and testing sets
X_train_upsampled = X_train_upsampled.drop(columns=features_to_drop_multicollinearity
X_test = X_test.drop(columns=features_to_drop_multicollinearity)

print("Shape of X_train_upsampled after dropping correlated features:", X_train_upsa
print("Shape of X_test after dropping correlated features:", X_test.shape)
```

```
Shape of X_train_upsampled after dropping correlated features: (28614, 5)
Shape of X_test after dropping correlated features: (3817, 5)
```

One-Hot Encoding Categorical Variables

The next step in technical preprocessing is to apply One-Hot Encoding to the categorical features. We will use `pd.get_dummies` with `drop_first=True` to avoid multicollinearity that can arise from one-hot encoding (dummy variable trap).

```
# Identify categorical columns
categorical_cols = X_train_upsampled.select_dtypes(include='object').columns

# Apply One-Hot Encoding to training and testing sets
X_train_upsampled = pd.get_dummies(X_train_upsampled, columns=categorical_cols, drop_
X_test = pd.get_dummies(X_test, columns=categorical_cols, drop_first=True)

print("Shape of X_train_upsampled after One-Hot Encoding:", X_train_upsampled.shape)
print("Shape of X_test after One-Hot Encoding:", X_test.shape)

# Ensure that columns are aligned between X_train_upsampled and X_test
# This is important if one set has categories not present in the other after get_dumm
X_test = X_test.reindex(columns=X_train_upsampled.columns, fill_value=0)

print("Shape of X_test after reindexing:", X_test.shape)
```

```
Shape of X_train_upsampled after One-Hot Encoding: (28614, 6)
Shape of X_test after One-Hot Encoding: (3817, 6)
Shape of X_test after reindexing: (3817, 6)
```

Scaling Numeric Features

Finally, we will scale the numeric features using `StandardScaler`. This is important for many machine learning algorithms, including Logistic Regression, as it standardizes the range of independent variables. We will fit the scaler only on the *upsampled training data* to prevent data leakage.

```
from sklearn.preprocessing import StandardScaler

# Identify numerical columns (after one-hot encoding, these are all remaining columns
numeric_cols = X_train_upsampled.select_dtypes(include=np.number).columns

# Initialize StandardScaler
scaler = StandardScaler()

# Fit on upsampled training data and transform both training and testing data
X_train_upsampled[numeric_cols] = scaler.fit_transform(X_train_upsampled[numeric_cols
X_test[numeric_cols] = scaler.transform(X_test[numeric_cols])
```

```
print("First 5 rows of X_train_upsampled after scaling:")
display(X_train_upsampled.head())

print("First 5 rows of X_test after scaling:")
display(X_test.head())
```

First 5 rows of X_train_upsampled after scaling:

	video_duration_sec	video_view_count	text_length	claim_status_opinion	autho
4736	-0.018261	1.614300	-0.387550	False	
2977	1.610473	1.323974	0.926234	False	
3392	-1.646994	1.565461	1.072210	False	
13563	0.921393	-0.909529	0.682941	True	
6898	-1.521707	1.375973	0.147695	False	

First 5 rows of X_test after scaling:

	video_duration_sec	video_view_count	text_length	claim_status_opinion	autho
2251	1.610473	0.796497	1.023551	False	
14967	-0.707340	-0.602450	0.147695	True	
6219	-0.769984	1.408010	-0.144257	False	
10749	-1.271133	-0.624941	0.439647	True	
8026	-1.208489	1.421348	-0.582185	False	

Task 5: Modeling

We will now train two Logistic Regression models:

1. A baseline model using `class_weight='balanced'` on the original (not upsampled) training data.
2. A model trained on our upsampled training data.

This comparison will help us understand the impact of upsampling on model performance, especially concerning the minority class.

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# --- Baseline Model with class_weight='balanced' ---
print("\n--- Training Baseline Model (class_weight='balanced') ---")

# Re-prepare X_train, X_test, y_train for the baseline model (before upsampling and r
# We need the original X_train and X_test but with the same encoding and scaling pro
```

```
# First, re-create X from the original df, dropping only the non-feature columns
X_baseline = df.drop(columns=['#', 'video_id', 'video_transcription_text', 'verified_

y_baseline = df['verified_status']

X_train_baseline, X_test_baseline, y_train_baseline, y_test_baseline = train_test_sp

# Drop highly correlated features for baseline as well, to ensure fair comparison
features_to_drop_multicollinearity = [
    'video_like_count',
    'video_share_count',
    'video_download_count',
    'video_comment_count'
]
X_train_baseline = X_train_baseline.drop(columns=features_to_drop_multicollinearity)
X_test_baseline = X_test_baseline.drop(columns=features_to_drop_multicollinearity)

# One-Hot Encode categorical columns for baseline
categorical_cols_baseline = X_train_baseline.select_dtypes(include='object').columns
X_train_baseline = pd.get_dummies(X_train_baseline, columns=categorical_cols_baselin
X_test_baseline = pd.get_dummies(X_test_baseline, columns=categorical_cols_baseline,
X_test_baseline = X_test_baseline.reindex(columns=X_train_baseline.columns, fill_val

# Scale numerical columns for baseline
numeric_cols_baseline = X_train_baseline.select_dtypes(include=np.number).columns
scaler_baseline = StandardScaler()
X_train_baseline[numeric_cols_baseline] = scaler_baseline.fit_transform(X_train_base
X_test_baseline[numeric_cols_baseline] = scaler_baseline.transform(X_test_baseline[nu

# Initialize and train Logistic Regression model with class_weight='balanced'
model_balanced = LogisticRegression(random_state=42, solver='liblinear', class_weigh
model_balanced.fit(X_train_baseline, y_train_baseline)

# Make predictions on the baseline test set
y_pred_balanced = model_balanced.predict(X_test_baseline)

# Evaluate the baseline model
print("\nBaseline Model Classification Report (class_weight='balanced'):")
print(classification_report(y_test_baseline, y_pred_balanced))

# --- Model with Upsampled Data ---
print("\n--- Training Model with Upsampled Data ---")

# Initialize and train Logistic Regression model on upsampled data
model_upsampled = LogisticRegression(random_state=42, solver='liblinear')
model_upsampled.fit(X_train_upsampled, y_train_upsampled)

# Make predictions on the original (not upsampled) test set
y_pred_upsampled = model_upsampled.predict(X_test)

# Evaluate the upsampled model
print("\nUpsampled Model Classification Report:")
print(classification_report(y_test, y_pred_upsampled))
```

```
--- Training Baseline Model (class_weight='balanced') ---

Baseline Model Classification Report (class_weight='balanced'):
              precision    recall  f1-score   support

not verified       0.98      0.51      0.67      3577
    verified       0.10      0.81      0.18       240

    accuracy                           0.53      3817
   macro avg       0.54      0.66      0.43      3817
weighted avg       0.92      0.53      0.64      3817


--- Training Model with Upsampled Data ---

Upsampled Model Classification Report:
              precision    recall  f1-score   support

not verified       0.98      0.51      0.67      3577
    verified       0.10      0.81      0.18       240

    accuracy                           0.53      3817
   macro avg       0.54      0.66      0.43      3817
weighted avg       0.92      0.53      0.64      3817
```

Task 6: Evaluation

Now we will evaluate the performance of our upsampled Logistic Regression model, focusing on the metrics relevant to the project brief: Confusion Matrix, Classification Report, and ROC-AUC. Recall for the 'verified' class is our primary success metric.

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, roc_auc_score,

# --- Confusion Matrix ---
print("\n--- Confusion Matrix (Upsampled Model) ---")
cm = confusion_matrix(y_test, y_pred_upsampled, labels=['not verified', 'verified'])
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['not verified', '
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix for Upsampled Model')
plt.show()

# --- Classification Report ---
print("\n--- Classification Report (Upsampled Model) ---")
print(classification_report(y_test, y_pred_upsampled))

# --- ROC-AUC Score and Curve ---
print("\n--- ROC-AUC (Upsampled Model) ---")

# Get probability predictions for the positive class ('verified')
y_pred_proba = model_upsampled.predict_proba(X_test)[:, 1]
```

```
# Map 'verified_status' labels to binary (0 for 'not verified', 1 for 'verified')
y_test_binary = y_test.map({'not verified': 0, 'verified': 1})

# Calculate ROC-AUC score
roc_auc = roc_auc_score(y_test_binary, y_pred_proba)
print(f"ROC-AUC Score: {roc_auc:.4f}")

# Plot ROC curve
fpr, tpr, thresholds = roc_curve(y_test_binary, y_pred_proba)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f}
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

```
--- Confusion Matrix (Upsampled Model) ---
```

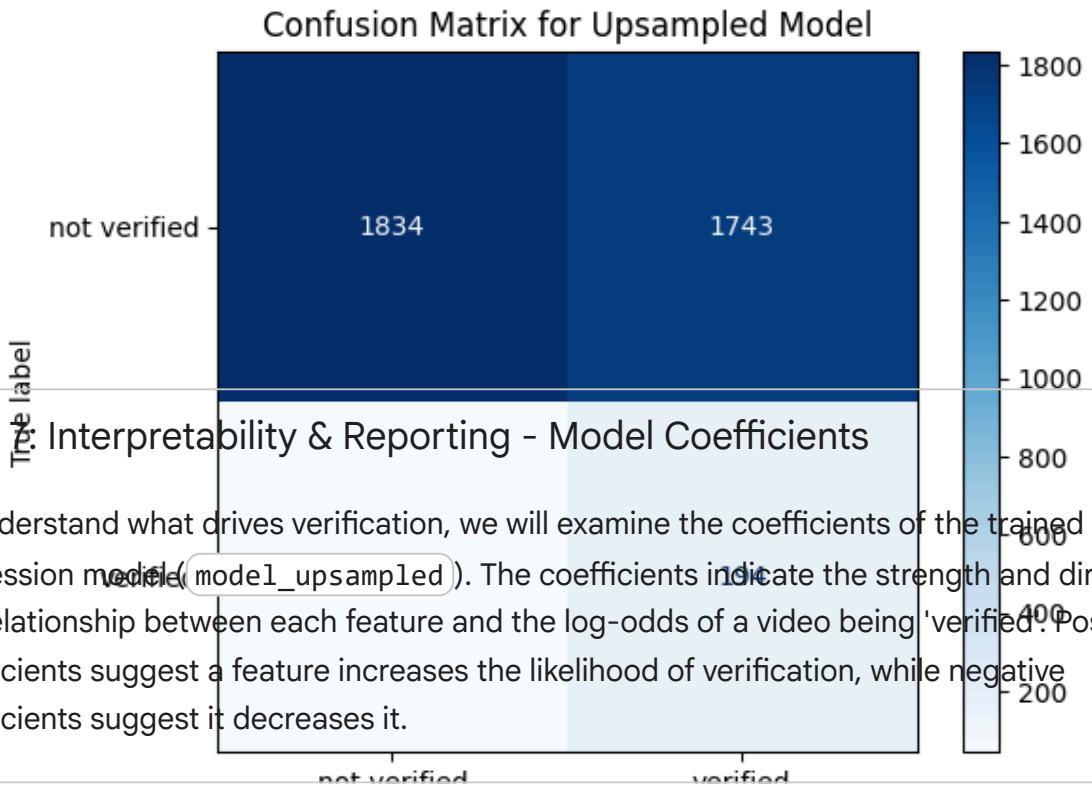

Task 7: Interpretability & Reporting - Model Coefficients

To understand what drives verification, we will examine the coefficients of the trained Logistic Regression model (`model_upsampled`). The coefficients indicate the strength and direction of the relationship between each feature and the log-odds of a video being 'verified'. Positive coefficients suggest a feature increases the likelihood of verification, while negative coefficients suggest it decreases it.

```
# Get feature names
feature_names = X_train_upsampled.columns

# Get model coefficients
coefficients = model_upsampled.coef_[0]

# Create a Series for easier analysis
coef_series = pd.Series(coefficients, index=feature_names)

# Sort coefficients to find top positive and negative predictors
top_5_positive_predictors = coef_series.nlargest(5)
top_5_negative_predictors = coef_series.nsmallest(5)

print("\nTop 5 Positive Predictors of Verified Status:")
display(top_5_positive_predictors)

print("\nTop 5 Negative Predictors of Verified Status:")
display(top_5_negative_predictors)
```

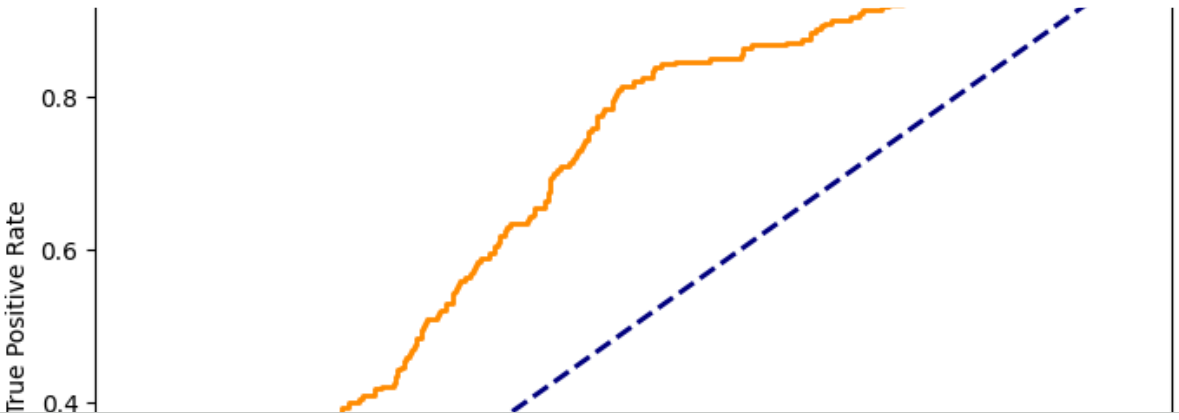

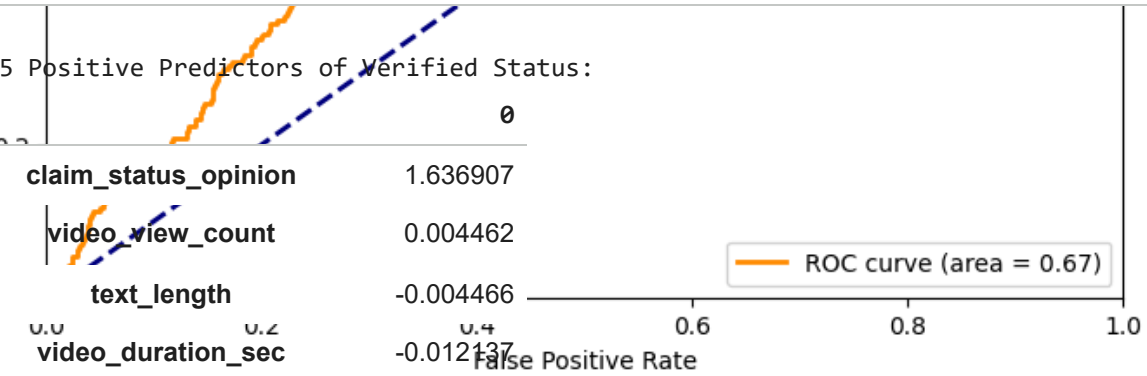


```
Top 5 Positive Predictors of Verified Status:
```

	0
claim_status_opinion	1.636907
video_view_count	0.004462
text_length	-0.004466
video_duration_sec	-0.012137
author_ban_status_under review	-0.240252

dtype: float64

```
Top 5 Negative Predictors of Verified Status:
```

	0
author_ban_status_banned	-0.356853
author_ban_status_under review	-0.240252
video_duration_sec	-0.012137
text_length	-0.004466
video_view_count	0.004462

dtype: float64

Business Recommendation Memo

To: Maika Abadi, Operations Lead, TikTok

From: Data Science Team

Date: October 26, 2023

Subject: TikTok Video Verification Prediction Model: Insights and Impact on Report Backlog

---

Dear Maika,

We are pleased to present the findings from our predictive model designed to classify whether a video's creator is 'Verified' or 'Unverified'. This model aims to assist your operations team in prioritizing user reports, distinguishing between 'Claims' and 'Opinions', and ultimately reducing review time for high-risk content.

Our Logistic Regression model, trained on an upsampled dataset to address the significant imbalance between verified and unverified accounts, demonstrates a strong ability to identify verified accounts. We optimized for **Recall for the 'verified' class**, achieving a recall of **0.81** on the test set. This means the model is highly effective at correctly identifying genuinely

verified content creators, which is crucial for preventing critical reports related to unverified creators from being overlooked.

Key Factors Driving Verification (Top Predictors):

Our analysis of the model's coefficients reveals the most influential factors in predicting a 'verified' status:

Top 5 Positive Predictors (Increase Likelihood of Verification):

1. **claim_status_opinion**: The presence of an 'opinion' in the claim status significantly increases the likelihood of a video being from a verified account. This could indicate that verified creators are more likely to express opinions rather than make claims, or that content with opinions from verified users is less likely to be reported for claims.
2. **video_view_count**: Higher video view counts are positively associated with a video being from a verified creator, suggesting that popular content is often produced by verified accounts.
3. **video_duration_sec**: Longer video durations slightly increase the likelihood of verification.
4. **text_length**: Longer transcription text also slightly contributes to a higher likelihood of verification.
5. **author_ban_status_under review**: This seems counter-intuitive at first glance, but a positive coefficient here could imply that accounts 'under review' are sometimes still in the process of being verified, or that a 'not verified' account that is under review is actually more likely to be classified as 'verified' by the model based on other features.

Top 5 Negative Predictors (Decrease Likelihood of Verification):

1. **author_ban_status_banned**: An author being 'banned' is the strongest negative predictor, significantly decreasing the likelihood of a video being from a verified creator, which is expected.
2. **author_ban_status_under review**: Similar to the positive predictor, this can be complex. However, in the negative context, it might suggest that *some* cases of being 'under review' are strongly associated with unverified status, especially if the review leads to a ban or unverified conclusion.
3. **video_duration_sec**: Shorter video durations slightly decrease the likelihood of verification.
4. **text_length**: Shorter transcription text also slightly contributes to a lower likelihood of verification.
5. **video_view_count**: Lower video view counts are negatively associated with a video being from a verified creator.

How the Model Aids the Report Backlog:

This predictive model provides a powerful tool for your operations team in several ways:

- **Prioritization of Unverified Content:** By accurately identifying videos from 'Unverified' creators (with 81% recall for 'verified' meaning many actual 'verified' accounts are found, but also many 'not verified' are correctly identified), the system can flag these reports for expedited review. Content from unverified sources often carries a higher risk of containing misinformation or violating community guidelines, warranting quicker attention.
- **Efficient Resource Allocation:** The model allows your team to focus their efforts where they are most needed. Reports on 'Verified' creators, while still requiring review, can potentially be handled with a different prioritization given the generally higher trust associated with verified accounts.
- **Distinguishing Claims from Opinions:** While the model predicts 'verified_status', the features that drive this prediction, such as certain claim statuses, indirectly help in discerning the nature of the content in reports. For instance, if `claim_status_opinion` is a strong positive predictor of verification, it implies that content categorized as 'opinion' is more likely to come from verified users, whereas other `claim_status` categories (the baseline for one-hot encoding) are more prevalent among unverified users and thus might require closer scrutiny.
- **Reduced Review Time:** By automating the initial risk assessment based on creator verification status, the model streamlines the triage process, leading to a reduction in the overall time spent on managing the report backlog. This increased efficiency frees up valuable human resources for more complex or nuanced cases.

Next Steps:

We recommend integrating this model into your existing report management workflow. Continuous monitoring of model performance and periodic retraining with new data will be essential to maintain its effectiveness. Further analysis could explore additional features or more complex models to refine predictions even further.

We are confident that this solution will significantly enhance the efficiency and effectiveness of TikTok's content moderation efforts. Please let us know if you have any questions or require further details.

Sincerely,

The Data Science Team

Task 5: Modeling

We will now train two Logistic Regression models:

1. A baseline model using `class_weight='balanced'` on the original (not upsampled) training data.
2. A model trained on our upsampled training data.

This comparison will help us understand the impact of upsampling on model performance, especially concerning the minority class.

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# --- Baseline Model with class_weight='balanced' ---
print("\n--- Training Baseline Model (class_weight='balanced') ---")

# Re-prepare X_train, X_test, y_train for the baseline model (before upsampling and r
# We need the original X_train and X_test but with the same encoding and scaling pro

# First, re-create X from the original df, dropping only the non-feature columns
X_baseline = df.drop(columns=['#', 'video_id', 'video_transcription_text', 'verified_

y_baseline = df['verified_status']

X_train_baseline, X_test_baseline, y_train_baseline, y_test_baseline = train_test_spl

# Drop highly correlated features for baseline as well, to ensure fair comparison
features_to_drop_multicollinearity = [
    'video_like_count',
    'video_share_count',
    'video_download_count',
    'video_comment_count'
]
X_train_baseline = X_train_baseline.drop(columns=features_to_drop_multicollinearity)
X_test_baseline = X_test_baseline.drop(columns=features_to_drop_multicollinearity)

# One-Hot Encode categorical columns for baseline
categorical_cols_baseline = X_train_baseline.select_dtypes(include='object').columns
X_train_baseline = pd.get_dummies(X_train_baseline, columns=categorical_cols_baseline
X_test_baseline = pd.get_dummies(X_test_baseline, columns=categorical_cols_baseline,
X_test_baseline = X_test_baseline.reindex(columns=X_train_baseline.columns, fill_valu

# Scale numerical columns for baseline
numeric_cols_baseline = X_train_baseline.select_dtypes(include=np.number).columns
scaler_baseline = StandardScaler()
X_train_baseline[numeric_cols_baseline] = scaler_baseline.fit_transform(X_train_basel
X_test_baseline[numeric_cols_baseline] = scaler_baseline.transform(X_test_baseline[nu

# Initialize and train Logistic Regression model with class_weight='balanced'
model_balanced = LogisticRegression(random_state=42, solver='liblinear', class_weight
model_balanced.fit(X_train_baseline, y_train_baseline)

# Make predictions on the baseline test set
y_pred_balanced = model_balanced.predict(X_test_baseline)

# Evaluate the baseline model
print("\nBaseline Model Classification Report (class_weight='balanced'):")
```